

1ST WORKSHOP ON ARTIFICIAL INTELLIGENCE FOR BIOMEDICAL DATA - AIBIO



xSTAE:

Explaining Classifier Decisions through EEG Signal Style Transfer Autoencoding

**Natalia Koliou (presenter),
Maria Sierra, Christoforos Romesis,
Stasinos Konstantopoulos, and Luis Montesano**



Funded by
the European Union

25 Oct 2025

Overview

1. Introduction & Related Work

2. Problem Statement

3. Proposed Methodology

4. Experiments

5. Conclusion & Future Work

Introduction & Related Work

xSTAE: Explaining Classifier Decisions through EEG Signal Style Transfer Autoencoding

Introduction & Related Work

- Understanding EEG-based sleep stage classification is crucial for sleep research and healthcare.
- Traditional models struggle with **misclassifications**, especially for minority sleep stages.
- Existing work interprets time-series classifiers, either by integrating explainers into the model (e.g., XTF-CNN, DeepVix) or using post-hoc techniques to highlight which input patterns drive predictions (e.g., timeXplain, LIME).
- **Counterfactual explanations** have been successful in visual domains but are underexplored for time-series EEG.
- We address this gap by generating **counterfactual EEG examples** that show what a misclassified instance should have looked like to be correctly classified.

Problem Statement

xSTAE: Explaining Classifier Decisions through EEG Signal Style Transfer Autoencoding

Problem Statement

- Counterfactual explanations provide **instance-based insights** for time-series classifiers.
- The goal is to explain why a classifier predicts a certain label:
 - Dataset: $D = \{x^{(i)}\}_{i=1}^N$, each $x^{(i)} \in \mathbb{R}^d$.
 - Classifier: $C : \mathcal{X} \rightarrow \mathcal{Y}$, labels $\mathcal{Y} = \{1, \dots, n\}$.
 - For input x , predicted label $y = C(x)$.
- For misclassified instances x_m with $C(x_m) \neq y^*(x_m)$, find **minimal modification** x'_m such that $C(x'_m) = y^*(x_m)$ and x'_m is close to x_m .
- $\Delta x_m = x'_m - x_m$ reveals dominant patterns influencing the classifier's decision.

Proposed Methodology

xSTAE: Explaining Classifier Decisions through EEG Signal Style Transfer Autoencoding

Proposed Methodology

- **xSTAE**: A generative framework using class-conditional **autoencoders**.
- Each autoencoder $E_{\text{tgt}} : \mathcal{X} \rightarrow \mathcal{X}_{\text{tgt}}$ is trained to reconstruct any input while restyling it toward a target class $y_{\text{tgt}} \in \mathcal{Y}$.
- For input x with $y_{\text{tgt}} \neq C(x)$, the autoencoder generates a counterfactual:

$$x' = E_{\text{tgt}}(x) \quad \text{s.t.} \quad C(x') = y_{\text{tgt}}$$

- Comparing x and x' reveals the patterns in x responsible for the classifier's original decision.
- Training uses dual loss function:
 1. **Identity**: Ensures x' remains similar to x , using a distance function $d(x, x')$.
 2. **Classification**: Ensures x' is classified as the target, by comparing $C(x')$ with y_{tgt} .

Experiments

xSTAE: Explaining Classifier Decisions through EEG Signal Style Transfer Autoencoding

Setup

We use the **BOAS EEG dataset**, with 128 full-night recordings from 2 channels (256 Hz), to explain classifier errors across 4 sleep stages (N1, N2, N3, REM).

- **EEG Data:**
 - 30s epochs ($n_s = 7680$ samples, $n_c = 2$).
 - FFT → 0.4–30 Hz filtering → segment into $n'_s = 300$ spectral slices.
 - Extract 3 features per slice (frequency, phase, amplitude).
- **Classifier:** Two-stage convolutional network over sequences of $k = 5$ epochs:

$$X_t^f = [e_{t-k+1}, \dots, e_t] \rightarrow \hat{y}_t$$

- **Autoencoders:** Hybrid networks with convolutional and attention layers:

$$X_f^t = e_t \oplus \text{pos} \longrightarrow \hat{X}_f^t \text{ (reconstructed input)}$$

Quantitative Results

- We pass the test dataset through each class-specific autoencoder, then classify the restyled outputs using the original classifier.
- High **classification accuracy** on restyled signals indicates the autoencoders have learned the patterns that the classifier uses for each sleep stage.

Metric	N1	N2	N3	REM
Accuracy	0.9997	0.9998	0.9260	0.9986

Qualitative Discussion

- **Example case:** Original EEG segment was labeled as N2 but misclassified as N1.
- xSTAE restyles the segment toward the correct class (N2), producing a counterfactual.
- Comparison reveals why the classifier erred:
 - Original signal lacked sufficiently prominent spikes.
 - Restyled signal emphasizes features around timesteps 55-60, 70-75, 90-95.
 - Classifier "expects" these more pronounced spikes to identify N2.
- Misclassification occurs not from wrong feature detection, but from **insufficient weighting** of existing patterns.

Conclusion & Future Work

xSTAE: Explaining Classifier Decisions through EEG Signal Style Transfer Autoencoding

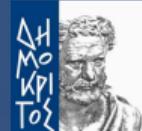
Conclusion

- **xSTAE** explains classifier errors by restyling misclassified EEG instances into correctly classified ones.
- Autoencoders balance **identity loss** (keep instance similar) and **classification loss** (push to correct label), revealing features the classifier missed.
- Contributions:
 - Ground counterfactual explanations in time-series EEG classification.
 - Identify spectral representations suitable for EEG signals.
 - Validate on open BOAS dataset and release full experimental setup.

Future Work

- Explore **alternative identity losses** to highlight meaningful changes (e.g., bigger local changes, selective brainwave bands).
- Conduct **expert trials** to refine interpretability of restyled EEGs.
- Investigate linking insights from misclassifications to **actionable guidance** at the data or confidence level, keeping xSTAE model-agnostic.

1ST WORKSHOP ON ARTIFICIAL INTELLIGENCE FOR BIOMEDICAL DATA - AIBIO



Thank you for your attention!

**Natalia Koliou (presenter),
Maria Sierra, Christoforos Romesis,
Stasinos Konstantopoulos, and Luis Montesano**



Funded by
the European Union

25 Oct 2025